

Contato

muraleob@gmail.com

www.linkedin.com/in/leonardo-murakami (LinkedIn)

Principais competências

Karpenter

Amazon SQS

Kong API Gateway

Languages

English (Native or Bilingual)

Portuguese (Native or Bilingual)

Spanish (Elementary)

Certifications

Python Programmer

Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning

Convolutional Neural Networks in TensorFlow

Data Analyst with Python

Data Scientist with Python

Honors-Awards

Menção Honrosa na OBMEP 2018

Leonardo Murakami

SRE @ Loft

São Caetano do Sul, São Paulo, Brasil

Resumo

Estudante de engenharia com forte interesse no campo da Ciência de Dados e Inteligência Artificial. Interessado em trabalhar na vanguarda do desenvolvimento de Inteligencia Artificial e Aprendizado de Maquina

Experiência

Loft

6 anos 10 meses

Mid-Level Site Reliability Engineer

março de 2024 - Present (2 anos 7 meses)

São Paulo, Brasil

Operating at the intersection of Site Reliability Engineering (SRE), Platform Engineering, DevEx, and LLMOps. Responsible for driving multi-cluster Kubernetes standardization, building self-service Internal Developer Platforms (IDP), enforcing GitOps/IaC governance, and optimizing cloud infrastructure costs across AWS ecosystems.

Key Achievements & Responsibilities:

- Internal Developer Platform (IDP) & DevEx: Engineered an automated self-service application bootstrap engine using GitHub Actions, Terraform, Terragrunt, and HashiCorp Vault. Reduced application onboarding and infrastructure provisioning time by 90% (from 4 hours to ~20 minutes) across SPA, Kotlin, and microservice architectures.
- Kubernetes Architecture & Reorganization (EKS): Leading the multi-phase migration and architectural standardization of 24 EKS Kubernetes clusters to enhance environment isolation, multi-level security governance, and granular FinOps cost tracking.
- AI Infrastructure & LLMOps (PoC): Architected and benchmarked a self-hosted open-weight LLM inference platform for Agentic Coding on EKS using NVIDIA GPUs, Karpenter (autoscaling), Knative, KServe, vLLM, and

Bifrost. Evaluated serverless scale-to-zero capabilities to drastically lower LLM inference unit economics.

- **API Gateway & Traffic Governance:** Governed API traffic infrastructure using Kong API Gateway, conducting architecture reviews, setting up external routing rules, and configuring security middleware (JWT authentication plugins).
- **FinOps & Cloud Cost Optimization:** Achieved a ~10% global cost reduction in Datadog by deprecating default agents in non-production environments and establishing an opt-in model for selective log and metric ingestion.
- **Incident Management & Reliability:** Provided tier-3 incident response and on-call support, debugging complex distributed systems issues spanning Kubernetes, PostgreSQL (RDS) query bottlenecks, AWS IAM permissions, and SQS/Datadog pipelines.

Junior Site Reliability Engineer

março de 2023 - março de 2024 (1 ano 1 mês)

São Paulo, Brasil

Focused on cloud infrastructure maintenance, continuous delivery automation, Kubernetes workload management, and operational toil reduction for product engineering teams.

Key Achievements & Responsibilities:

- **Legacy Migration (EC2/Jenkins to EKS):** Led the technical migration of development workloads in the Finance domain from legacy EC2 instances and Jenkins pipelines to modern EKS Kubernetes clusters using GitOps (Argo CD).
- **Kubernetes Workload Operations:** Maintained and scaled containerized workloads using Argo CD, Helm charts, and GitHub Actions CI/CD pipelines.
- **Toil Reduction & Automation:** Built internal scripts and CLI automations to streamline cloud resource provisioning, significantly reducing manual operational toil for engineering squads.
- **Cross-Squad Technical Support:** Acted as an embedded technical consultant for product squads, resolving issues around observability (Datadog), cloud permissions (AWS IAM), API routing (Kong), and database connectivity.

Machine Learning Engineer

maio de 2022 - março de 2023 (11 meses)

São Paulo, Brasil

Managed MLOps lifecycles, automated ML training infrastructure, and productized automated valuation models (AVM) for real estate pricing at scale.

Key Achievements & Responsibilities:

- MLOps & Pipeline Sustainance: Governed the end-to-end MLOps lifecycle for production Machine Learning models, ensuring low-latency inference, model drift monitoring, and high system availability.
- ML Training Infrastructure: Provisioned and managed scalable AWS SageMaker infrastructure, enabling Data Scientists to accelerate experiment iteration and model training cycles.
- Model Productization: Supported the production deployment and API integration of proprietary Automated Valuation Models (AVM) into Loft's core microservices architecture.

Junior Data Scientist

junho de 2021 - julho de 2022 (1 ano 2 meses)

São Paulo, Brasil

Developed statistical and machine learning models for real estate valuation, comparative market analysis, and automated pricing algorithms (AVM).

Key Achievements & Responsibilities:

- Automated Valuation Models (AVM): Designed, trained, and validated core real estate pricing algorithms (including production models Robson, Leia, Amadeu, and David) utilizing property registry data, market listings, and geo-spatial feature sets.
- Feature Engineering & Data Pipelines: Built robust data pipelines in Python and SQL to extract, transform, and engineer spatial and transactional features for pricing accuracy.
- Model Monitoring & Quality Control: Established statistical metrics and automated dashboards to track valuation accuracy, price variance, and model performance over time.

Data Science Intern

fevereiro de 2020 - junho de 2021 (1 ano 5 meses)

Transitioned from the Summer Internship program to a full-time internship, deepening my involvement in data-driven product decisions, automated valuation, and ML model development for the real estate market.

Key Achievements & Responsibilities:

- **Data Pipelines & Feature Extraction:** Built and optimized analytical pipelines in Python and SQL to extract, clean, and transform transactional and geo-spatial data from corporate databases.
- **Machine Learning Model Support:** Assisted senior Data Scientists in the validation, testing, and monitoring of early-stage predictive models used for real estate valuation.
- **Business Intelligence & Reporting:** Created automated dashboards and conducted exploratory data analysis (EDA) to support strategic decision-making across product and business teams.

Data Science Summer Intern

dezembro de 2019 - fevereiro de 2020 (3 meses)

Selected for Loft's competitive Summer Internship program, where I was introduced to the intersection of real estate technology and data science.

Key Achievements & Responsibilities:

- **Exploratory Analysis:** Conducted initial exploratory data analysis (EDA) on large-scale property datasets to identify market trends and data anomalies.
- **Experiment Design:** Supported the design and setup of A/B tests and early predictive experiments to evaluate the impact of various real estate features on pricing.
- **Cross-functional Collaboration:** Gained foundational experience working within an agile tech environment, collaborating closely with data engineering and product teams.

Grupo Turing

Machine Learning Researcher

setembro de 2019 - março de 2020 (7 meses)

São Paulo Area, Brazil

Formação acadêmica

Instituto de Matemática e Estatística - Universidade de São Paulo
(IME-USP)

Bacharelado, Estatística · (2021 - 2025)

Escola Politécnica da USP

Bachelor of Engineering - BE, Electrical Engineering - Habilitation: Computer
Engineering · (2020 - 2021)

Escola Politécnica da USP

Bachelor of Engineering - BE, Mining and Mineral Engineering · (2019 - 2020)